

Searching heterogeneous data-sources: Master thesis problem statement

Vidmantas Zemleris, October 18, 2012

1 Introduction

At large scientific collaborations like the CMS Experiment at CERN's LHC that includes more than 3000 collaborators data usually resides on a fair number of autonomous and heterogeneous proprietary systems each serving it's own purpose ¹. As data stored on one system may be related to data residing on other systems², users are in need of a centralized and easy-to-use solution for locating and combining data from all these multiple services.

Using a highly structured language like SQL is problematic because users need to know not only the language but also where to find the information and also lots of technical details like schema. A data integration system based on simple structured queries is already in place. Various improvements including support for less restricted keyword queries, improvements to system's usability and performance still have to be researched.

2 Case study: the CMS Experiment at CERN

Users' information need may vary greatly depending on their role, however most of the time they are interested in locating a *full set* of entities matching some selection criteria, e.g.:

- find all *files* from *dataset(s)* matching wild-card query each containing some of the 'interesting' *runs* from a list provided (Release validation teams)
- find (all) *datasets* related some specific physics phenomena³ together with conditions describing how this data was recorded by detector or simulated which are present in separate autonomous system than the datasets (Physicists)
- find all *datasets* matching some pattern stored at a given *site* (filtering attributes from separate services)

For more use-cases of data retrieval at CMS Experiment see [DGK⁺08].

The Data Aggregation System

The Data Aggregation System (DAS)[KEM10, BKEM11] was created which allows integrated access to a number of proprietary data-sources by processing user's queries on demand - it determines data-sources are able to answer⁴, queries them, merges the results and caches them for subsequent use. DAS uses *Boolean retrieval model* as users are often interested in retrieving ALL the items matching their query.

Currently the queries specify what entity the user is interested in (dataset, file, etc) and provide selection criteria (attribute=value, name BETWEEN [v1, v2]) operators. The combined query results could be later 'piped' for further filtering, sorting or aggregation (min, max, avg, sum, count, median), e.g.:

```
dataset=*RelVal* | grep dataset.nevents >1000 | avg(dataset.size), median(dataset.size)
```

The query above would return average and median datasets sizes of ones containing *RelVal* in their name having more than 1000 events.

Queries could be run either from web browser or through command line interface where the results could be fed into another application (e.g. program doing physics analysis or automatic software release validation).

¹For instance, at CERN, due to many reasons (e.g. research and need of freedom, politics of institutes involved) software projects usually evolve in independent fashion resulting in fair number of proprietary systems[KPGM00]. Further high turnover makes it harder extending these systems

²For example, datasets containing physics events are registered at DBS, while the physical location of files is tracked by Phedex which also takes care of their transfers within the worldwide grid storage

³In case of dataset this data is present in filename or run

⁴This is done by a mapping between flat mediated schema ('DAS keys') into web-service methods and their arguments. Then system queries all services that could provide a result of expected type with given parameters

3 Problem statement

3.1 User Interface and ease of use

For an IR system with wide variety of users, it is important to provide an easy to use interface with fairly flat learning curve, while at the same time not losing support of fairly complex queries. Even a simple structured query language and entity names over the mediated schema at first may seem hard to learn⁵. Therefore, it shall be useful to guide new users interactively through the process of building of their query. Supporting non-structured keyword queries is also worth investigation as quite many users reported they are missing this Google-like search experience.

3.2 Supporting complex queries

Currently DAS can only process queries that could be directly mapped to data providers' APIs, but not the queries where results of one API need to be fed into another APIs⁶. Further, to be flexible DAS only keeps a list of 'mediated entities' and how they could be retrieved, but do not enforce any predefined schema (i.e. entity fields/attributes) nor exposes it⁷ before querying for some particular entity(-ies). This makes it fairly hard to execute complex queries without a priori knowledge. User has to know what queries and how have to be combined (depends on the mediated schema and what APIs are available), and has either to combine the results by hand or to write a program does that. Some of the questions to be answered:

- could we select meaningful and sufficiently efficient composition of services even using less detailed source descriptions than those used in Information Manifold (see Literature Review)?
- could we use the current descriptions that only define entity the result and API parameters?

3.3 Performance

DAS is based on *Virtual Integration* where data is left at the sources allowing data to be volatile (e.g. new records gathered or existing records updated; service descriptions may be changed; new services added), while some subsets of the data could be fairly static (e.g. datasets stored at DBS do not change often). Further as the total data that is stored at data-providers could be fairly large (order of 1TB per year) and as their main goal is supporting production (data taking from the CMS detector), the providers may impose some constraints (e.g. only certain queries allowed). Therefore, it is important to balance between returning locally cached query results quickly and between issuing queries to slow data-services to get fresh results. Issues to be addressed includes:

- Low data providers' performance (e.g. on DBS some queries requires joining many very large tables)
- At the moment all queries are put into one pool and has to wait until some threads are available. If a couple of heavy queries were submitted earlier, even light queries would have to wait long. Explore more advanced Query prioritization (e.g. we could have some query cost model based on history or predefined scores per API)
- Currently result items are cached for a fairly short period of time (5min-1h) and then completely discarded, however many entities are not changing that often - Explore more intelligent caching
- Since DAS does aggregation across multiple data providers, given their current APIs (with no ordering and limit functionality), DAS has to fetch ALL records matching the query instead of only the first page. Although it would be possible to do partial results for data coming from a single data-source (while still loading in background, or the providers (or us) would have to modify their APIs).
- Efficient distributed search: as filtering criteria may reside on two or more autonomous sources, given the current APIs much more items than in the final result may need to be fetched from each source. Items that do not meet all selection criteria could be filtered out only afterwards in DAS. To improve the performance, data provider's APIs for the most used queries may need to be redesigned.

⁵on the other hand, at CERN, the names in the mediated schema are refer to entities that in the real-world are fairly consistently named (though there may exist slight differences in their naming, whereas DAS is to provide the 'golden' naming convention).

⁶except a couple of manually hard-coded "views" in a virtual "combined" data provider

⁷Actually the fields of each API (an entity could be mapped to a number of these) could be inferred from past queries or by predefined 'bootstrap' queries. Currently the field lists may differ slightly depending on the parameters

3.4 Generic connector accessing relational databases

Generic connector able to access relational databases with minimal manual integration could be useful for integrating proprietary systems that were not yet integrated and there's no resources available to build APIs. A possible use-case could be the *Prep* database, but this still has to be discussed. Depending on the needs it could be either of exploratory or API/Query Forms based approach, where fake APIs are defined in terms of SQL queries (or even better in a simplified form that would generate SQL by taking schema into account) There was a project⁸ at CERN which tried to generalize DBS-Query Language for any database.

4 Proposed solutions

4.1 Ease of use

- Interactive guidance on how to compose queries to ease the learning (could also include some examples of popular queries):
 1. What entity are you searching for?
 2. How could you identify it (i.e. what do you know)?
 3. What do you want to do with these results (after seeing the results)? Possible options: list items; select only specific columns; do aggregation; use in command line/another program
- Given a keyword query, suggest the best matching structured queries (as a ranked list)⁹:
 - the algorithm will generate mappings between each keyword to either an entity name, selection attribute name or its value (resembling 'Query forms' and 'Keyword cleaning' approaches)
 - ranking could be based on: how closely keywords are matching (some or all) required API's input parameters, popularity of certain queries and users feedback
 - build an inverted index of attribute values with some good full-text search toolkit¹⁰. If no match found (e.g. new entry not yet in our cache) pattern matching could be used to make a best guess
 - the system could improve its results from users' feedback (e.g. by HMM, will need to decide)
 - multi-wildcard queries (e.g. dataset=*Zmm*special*RECO*) are common at CMS that will need special attention¹¹.
 - knowing all possible APIs parameter values would be beneficial (but not yet sure if available¹²) and could be used to improve the keyword search (better mapping from keywords to API inputs) or even the structured search (e.g. query cleaning/term auto-completion)
 - * having direct access to DB, one could incrementally index target database¹³ tables mapping it's columns into API parameters (manually or automatically).
- (?) JavaScript-based query interpreter to ease query writing: could suggest search attribute and entity attribute names. also taking into account some possible ambiguous namings (auto completion)

⁸<https://github.com/vkuznet/PyQueryBuilder>

⁹As DAS uses Boolean IR model we could only rank specific structured queries, but not the results itself.

¹⁰e.g. Xapian (<http://xapian.org/>) implemented in C++ which that is known for its flexibility, or Sphinx (<http://sphinxsearch.com/>)

¹¹Matching to structured query will be harder and guessing even worse. Also for performance reasons some web-services will introduce limitations on the format of wildcard queries, while DAS wants still to be flexible at least for datasets

¹²problem: the owners of services may not want to provide direct access to DB because of no trust, security or performance issues. Otherwise, specific APIs could be developed to return possible parameter values but this is quite an overhead

¹³E.g. an inverted index of values in Oracle/MySQL DB tables can be built to keep the list of possible terms. For instance, Sphinx full text search engine can access DBs directly and also supports incremental indexing but it needs a bit of manual configuration <http://sphinxsearch.com/docs/current.html#delta-updates>

4.2 Supporting complex queries

As most entities in the mediated schema (27 of them) are already mapped to multiple services (see the mapping here: <https://cmsweb.cern.ch/das/services>) in theory these could be taken as a graph could to build (some of) the more complex queries. *At the moment this still needs some further consideration/planning, as the source descriptions we have are not so detailed as for example *Information Manifold*'s.*

4.3 Performance

More intelligent caching

Some of the data entities instances changes very rarely. For example in DBS system old datasets would never change, while new entities are constantly added (still some of their attributes may change, e.g. validity of a dataset). Cached copy may be shown by default while up to date results could be retrieved on user's request. An automatic change rate prediction could be useful to efficiently balance between caching and retrieving results. On the other hand some result items may not make sense to be cached, as there may be too many of valid input parameter combinations (e.g. now there are 900M of (*run*, *file*) combinations).

Also consider that:

- we could also have different validity dates for certain fields. if no volatile fields are not explicitly requested, even a very old cache could be used.
- pre-fetching common queries: determining manually and/or automatically

Data providers performance issues

In the cases when new records are coming but the existing ones are not changing much, continuous view maintenance that computes only differences from earlier results could be a good solution to improve the performance especially for often used queries.

Use either *materialized refresh fast views with query rewriting* (Oracle; completely transparent for proprietary apps)[ea11] or some other continuous view maintenance tool (e.g. DBToaster¹⁴) to improve performance of heavy queries containing joins and/or aggregations.

This is exactly the case for the most popular expensive query (see query #7 in A.2 DAS log analysis) over DBS system (80GB + 280GB indexes): 'find files where run in [r1, r2, r3] and dataset=X' that needs to join all the biggest tables in that database.

Dataset (164K rows) -> Block (2M) -> Files (31M) -> FileRunLumi (902M) <- Runs (65K)

Having a materialized view with all these tables joined together would allow answering such queries much quicker. Given low change rates (in comparison to data already present), maintaining the view should be also comparatively cheap with only expense of just couple of times of space.

Integrating distributed information efficiently

Bloom-join (which could be quite transparent and implemented even on DB side [pure SQL is possible for MySQL, to check for Oracle]- take a query and bit-vector as parameter) , lazy pagination (and order required for aggregation) - this is not yet supported by any of the data-service APIs

¹⁴<http://www.dbtoaster.org> that is being developed at EPFL

5 Preliminary Literature review

5.1 Overview

Since the late 1990's, several Enterprise Information Integration¹⁵ (EII) products have appeared in the market (e.g. *Information Manifold* by AT&T Lab) and an significant experience has been accumulated on: data integration formalisms, ways of describing of heterogeneous data sources and their abilities (e.g. RDBMS vs web form), query optimization (combining sources efficiently, source overlap, data quality, etc)[HAB⁺05]. Recent research in data integration mostly focused on methods of automatic, uncertainty and feedback driven (e.g. though statistical mediator) semi-automatic pay-as-you-go approaches[AD12, ch.19]).

The problem of keyword search over structured sources received significant attention within the last decade. Keyword search over relational, XML and other databases was explored from a number of perspectives: returning top-k ranked result tuples vs suggesting structured queries as SQL, performance optimization, user feedback mechanisms, as well as keyword searching over distributed sources. It is worth noting that when there is no need for 100% result exactness keyword search combined with probabilistic schema matching allows lightweight (pay-as-you-go) data integration with minimal human effort upfront[AD12, ch.16] but could improve its quality with time by using users' feedback.

In the setting of Boolean Keyword-based Retrieval over heterogeneous sources a keyword query could be translated into a ranked list of structured queries (similarly as SODA system that propose SQL, Query Forms that propose SQL templates or Keymantic trying to achieve this without accessing the data). In the extreme case of having no control over a web-service that do not publish its contents, techniques like Google's Deep-web crawl could index a subset of its contents enabling keyword search to some extent.

5.2 Composing multiple data sources efficiently: Query translation using Logical Views

The *Information Manifold*[LRO96] is virtual integration (EII) system that uses CARIN dialect of *description logics* (could be though as simplified datalog) for queries and for source descriptions. It describes each source as logical view over the global (mediated) schema, augmented with source capabilities (e.g. what are possible and the required input parameters for a view to do selections¹⁶). This allows designing algorithms that could allow efficiently answering complex queries that require composing multiple the data sources by finding maximally contained rewriting of (conjunctive) input query in terms of logical views provided by sources (that is, finding an optimal way to compose the sources).

For example (based on [Ull00]), consider such sources as expressed as views (on the left) in terms of global predicates of the mediated schema (on the right) in datalog notation:

```
v1(E, P, M) :- emp(E) & phone(E, P) & mgr(E, M). # employees, their phones and managers
v2(E, O, D) :- emp(E) & office(E, O) & dept(E, D). # offices and departments of employees
v3(E, P) :- emp(E) & phone(E, P) & dept(E, toy_dept). # phones of employees only in Toys dept.
```

Suppose we wanted to know Sally's phone and office. We express this by datalog in global predicates:

```
q1(P,O) :- phone(sally,P) & office(sally,O).
```

There are two minimal solutions (as the sources could be incomplete the full solution is union of the two):

```
answer1(P,O) :- v1(sally,P,M) & v2(sally,O,D).
answer2(P,O) :- v3(sally,P) & v2(sally,O,D).
```

Notice that the expansions of these solutions (e.g. `answer1_exp`) are not equivalent to q_1 , but the conjunctive queries that are closest and still contained in q_1 (as these are the only usable views at sources):

```
answer1_exp(P,O) :- emp(sally) & phone(sally,P) & mgr(sally,M) & emp(sally) & office(sally,O) & dept(sally,D). ■
```

After this *Information Manifold* would find an executable order that adheres the capabilities of the sources, by iteratively considering any sources whose input parameters are satisfied.

¹⁵Enterprise Information Integration (EII) is about 'integrating data from multiple sources *without* having to first load data into a central warehouse'[HAB⁺05, p.1]

¹⁶this makes these logical views quite similar to source APIs used by DAS, with difference that DAS currently only describe the parameters APIs and only partially the results (as an entity type)

5.3 Keyword search over a Relational Database

The problem of Keyword search over Relational Databases (or also semi-structured sources like XML) has received a significant attention by the research community over the last decade.

The basic approach would first build an inverted index on database tables (usually only text columns). Then after finding all occurrences of the keywords, would try to construct join paths (based on Foreign keys) that would unite tuples containing the keywords.

A number of problem variations exist: returning only the ranked Top-k results vs. returning ranked list of possible queries, while some systems would even allow generating more complex queries including aggregations, etc (SQAK, SODA[BJK⁺12]).

Ranking Query Templates based on keyword query

A simple way to access relational database could be through a set of predefined named query templates (SQL with selection parameters or operators still to be specified) exposed to a user as a Form that the user has to fill in.

[CBC⁺09] proposes alternative approach for processing keyword queries over relational databases: given a keyword query, instead of returning database tuples one could rank query forms that best matches the query for user to choose the right one (if they are properly named this is fairly easy). The ranking is based on checking matching of keywords to table names in templates and to column values.

An interesting feature of this approach is that a Query Template is functionally similar to any autonomous web service (which given the parameters would in turn execute that query on its database). In case of the Data Aggregation System, a user after entering a keyword query could be provided with a ranked list of structured queries (attribute=value) that could be processed given data source constraints (e.g. parameters required) and if needed he could refine his search (e.g. provide more parameters).

Keyword query cleaning

Keyword queries are often ambiguous, may contain misspellings or multiple keywords that refer to the same attribute value, therefore [PY08] suggested to perform query cleaning before proceeding to subsequent more computationally expensive steps (e.g. exploring all the possible join paths).

Further employing some machine learning method like HMM[Pu09] would allow to incorporate user's feedback (even the fact that user has chosen n-th result as a query to be executed is a good clue).

SODA: Meta-data approach

With a goal to bridge the increasing gap between high-level (conceptual, business) and low level (physical) representations of data, researchers from *ETHZ* have been investigating Generation of SQL for Business users over a very complex data warehouse at *Credit Suisse*. For converting natural language queries into SQL statements, in addition to what used by earlier approaches they used meta-data describing the schema at both physical, conceptual and logical levels extended with DBpedia (for synonyms, etc) and domain ontologies (to capture business concepts like 'wealthy customer').

Even on a large data-warehouse of 220GB data with a complex schema of 400+ tables they reported that if good meta-data is available, generating even fairly complex SQL (e.g. n-way joins with aggregations) is quite feasible for a computer. That would make it 'much easier for business users to interactively explore highly-complex data warehouses' [BJK⁺12, p.932]. The users also reported system's potential a) for analysing the schema and learning patterns about it and b) as tool to help documenting legacy systems.

5.4 Keyword search over web services

What if there is no access to index data terms

[BDG⁺10, BDG⁺12] explores the case then there is no possibility to index the data terms, e.g. then a DB is behind a wrapper (e.g. accessible only through a *Web form* in Hidden Web or a *web-service*) then crawling is generally not possible. In Keymantic[BDG⁺10] a 'keyword query is processed as follows: First, all keywords that correspond to metadata items (e.g., field names) are extracted. The remaining keywords are considered

as possible input fields. Second, the likelihood of a remaining keyword to a metadata item is computed in order to rank different options to execute the keyword query on the Hidden Web database'[BJK⁺12, p.942].

Deep web search at Google

There are two approaches to web scale search over deep-web (here Google mostly cares about web forms):

Deep-web surfacing - surface deep-web (e.g. web forms) adding their results into the standard search index easily allowing to using existing IR technology that scales well. There exist algorithms which allow to iteratively choose input parameters to the forms to surface a considerable part of the 'hidden' data without large overhead¹⁷.

Pay-as-you-go approach[MJC⁺07] - With this approach, 'a system starts with very few (or inaccurate) semantic mappings and these mappings are improved over time as deemed necessary'; there is NO single mediated schema over which users pose queries: queries are routed to to the relevant sources with help statistical methods that are used to model uncertainty at all levels: queries, mappings and underlying data.

TODO:
does - g
use PayC

¹⁷i.e. if choosing parameters in not smart way, a web form with just a couple of free inputs (or even dropdowns that's easier case), could yield as many results as a cross product of all input combinations.

6 Work status

- obtained DB copy of biggest data provider DBS (currently 80 GB + 200 GB indexes)
- preliminary literature review (to be continued deeper)
- initial analysis of DAS logs

Upcomming Work items:

- couple of fairly simple prototypes of UI/access patters for simpler DAS querying
- check performance improvements after creation of materialized view(s)
- interview more DAS users

References

- [AD12] Zachary Ives Anhui Doan, Alon Halevy. *Principles of data integration*. Number 9780124160446. Morgan Kaufmann, 2012. 497p.
- [BDG⁺10] Sonia Bergamaschi, Elton Domnori, Francesco Guerra, Mirko Orsini, Raquel Trillo Lado, and Yannis Velegrakis. Keymantic: semantic keyword-based searching in data integration systems. *Proc. VLDB Endow.*, 3(1-2):1637–1640, September 2010.
- [BDG⁺12] Sonia Bergamaschi, Elton Domnori, Francesco Guerra, Silvia Rota, Raquel Trillo Lado, and Yannis Velegrakis. Understanding the semantics of keyword queries on relational data without accessing the instance. In Roberto Virgilio, Francesco Guerra, Yannis Velegrakis, M. J. Carey, and S. Ceri, editors, *Semantic Search over the Web, Data-Centric Systems and Applications*, pages 131–158. Springer Berlin Heidelberg, 2012.
- [BJK⁺12] Lukas Blunschi, Claudio Jossen, Donald Kossmann, Magdalini Mori, and Kurt Stockinger. Soda: generating sql for business users. *Proc. VLDB Endow.*, 5(10):932–943, June 2012.
- [BKEM11] G Ball, V Kuznetsov, D Evans, and S Metson. Data aggregation system - a system for information retrieval on demand over relational and non-relational distributed data sources. *Journal of Physics: Conference Series*, 331(4):042029, 2011.
- [CBC⁺09] Eric Chu, Akanksha Baid, Xiaoyong Chai, AnHai Doan, and Jeffrey Naughton. Combining keyword search and forms for ad hoc querying of databases. In *Proceedings of the 2009 ACM SIGMOD International Conference on Management of data*, SIGMOD ’09, pages 349–360, New York, NY, USA, 2009. ACM.
- [DGK⁺08] A Dolgert, L Gibbons, V Kuznetsov, C D Jones, and D Riley. A multi-dimensional view on information retrieval of cms data. *Journal of Physics: Conference Series*, 119(7):072013, 2008.
- [ea11] P Lane et al. Oracle database data warehousing guide, 11g release 2 (11.2). chapter 9: Basic materialized views, September 2011.
- [HAB⁺05] Alon Y. Halevy, Naveen Ashish, Dina Bitton, Michael Carey, Denise Draper, Jeff Pollock, Arnon Rosenthal, and Vishal Sikka. Enterprise information integration: successes, challenges and controversies. In *Proceedings of the 2005 ACM SIGMOD international conference on Management of data*, SIGMOD ’05, pages 778–787, New York, NY, USA, 2005. ACM.
- [KEM10] Valentin Kuznetsov, Dave Evans, and Simon Metson. The cms data aggregation system. *Procedia Computer Science*, 1(1):1535 – 1543, 2010. [jce:titlejICCS 2010j/ce:titlej](#).
- [KPGM00] Christoph Koch, Paolo Petta, Jean-Marie Le Goff, and Richard McCatchey. On information integration in large scientific collaborations, 2000.
- [LRO96] A. Levy, A. Rajaraman, and J. Ordille. Querying heterogeneous information sources using source descriptions. Technical Report 1996-61, Stanford InfoLab, 1996.
- [MJC⁺07] Jayant Madhavan, Shawn R. Jeffery, Shirley Cohen, Xin (luna) Dong, David Ko, Cong Yu, Alon Halevy, and Google Inc. Web-scale data integration: You can only afford to pay as you go. 2007.
- [Pu09] Ken Q. Pu. Keyword query cleaning using hidden markov models. In *Proceedings of the First International Workshop on Keyword Search on Structured Data*, KEYS ’09, pages 27–32, New York, NY, USA, 2009. ACM.
- [PY08] Ken Q. Pu and Xiaohui Yu. Keyword query cleaning. *Proc. VLDB Endow.*, 1(1):909–920, August 2008.
- [Ull00] Jeffrey D. Ullman. Information integration using logical views. *Theoretical Computer Science*, 239(2):189 – 210, 2000.

A DAS Query logs (from 2011-06-21 to 2012-10-01)

From the logs it can be seen that queries requiring heavy joins are quite common (100K queries through command line interface) that makes it worth investigating the possible performance optimizations.

A.1 Common query patterns through Web browser

total valid queries: 569,408

not well formed queries (e.g. free text, typing mistakes, spam): 98,923

1	50.16%	(285605)	:	dataset dataset.name=?
2	13.54%	(77071)	:	site dataset.name=?
3	8.96%	(51035)	:	file dataset.name=?
4	5.88%	(33504)	:	run dataset.name=?
5	2.59%	(14739)	:	release dataset.name=?
6	2.11%	(12036)	:	config dataset.name=?
7	1.65%	(9420)	:	dataset run.run_number=?
8	1.34%	(7642)	:	block dataset.name=?
9	1.24%	(7084)	:	dataset site.name=?
10	1.15%	(6562)	:	dataset dataset.name=? release.name=?
11	1.10%	(6287)	:	parent dataset.name=?
12	0.96%	(5488)	:	file file.name=?
13	0.92%	(5245)	:	dataset dataset.name=? status.name=?
14	0.77%	(4363)	:	dataset release.name=?
15	0.75%	(4257)	:	file dataset.name=? run.run_number=?
16	0.68%	(3874)	:	run run.run_number=?
17	0.53%	(2994)	:	site site.name=?
18	0.45%	(2576)	:	site file.name=?
19	0.45%	(2556)	:	dataset file.name=?
20	0.43%	(2438)	:	lumi file.name=?
21	0.40%	(2282)	:	dataset dataset.name=? site.name=?
22	0.35%	(1999)	:	file block.name=?
23	0.35%	(1970)	:	dataset dataset.name=? run.run_number=?
24	0.29%	(1640)	:	group dataset.name=?
25	0.29%	(1631)	:	lumi run.run_number=?

Interesting non-valid queries:

- * keyword search: *herwig*/AODSIM

- * Users may like more complex combined queries:

- lumi file = (file dataset=/RelValProdTTbar/JobRobot-MC_42_V12_JobRobot-v1/GEN-SIM-RECO)■

- file,lumi dataset=/RelValProdTTbar/JobRobot-MC_42_V12_JobRobot-v1/GEN-SIM-RECO

- * Users mixing up the post and pre filters:

- file dataset=/MuEG/Run2011B-PromptReco-v1/AOD, file.size >1

- file dataset=/MinimumBias/Run2010A-vals-kim-v6/RAW-RECO* | grep run between [138923, 144086]■

A.2 Common query patterns through Command Line

total valid queries: 6,9M

non valid queries: 76K

1	39.56%	(2735728)	:	dataset dataset.name=?
2	22.80%	(1576886)	:	dataset dataset.name=? status.name=?
3	12.49%	(863636)	:	file dataset.name=?
4	12.44%	(860626)	:	run run.run_number=?
5	4.60%	(318248)	:	site dataset.name=?
6	2.18%	(150845)	:	run dataset.name=?
7	1.71%	(118404)	:	file dataset.name=? run.run_number=?
8	1.32%	(90988)	:	file block.name=?
9	1.26%	(87266)	:	file file.name=?
10	0.50%	(34482)	:	block site.name=?
11	0.35%	(24049)	:	lumi file.name=?
12	0.25%	(17282)	:	release dataset.name=?
13	0.18%	(12556)	:	parent file.name=?
14	0.09%	(6341)	:	file dataset.name=? lumi.number=? run.run_number=?
15	0.06%	(4306)	:	file dataset.name=? site.name=?
16	0.05%	(3547)	:	dataset dataset.name=? primary_dataset.name=? release.name=? tier.name=?
17	0.05%	(3352)	:	dataset file.name=?
18	0.01%	(996)	:	parent dataset.name=?
19	0.01%	(755)	:	dataset site.name=?

A Data providers statistics

(Some of the largest ones)

DBS:

DB size: 80GB + 200GB indexes, not many changes to existing (old) records

Largest tables:

Dataset (164K rows) -> Block (2M) -> Files (31M) -> FileRunLumi (902M) <- Runs (65K)■

Phedex: ~7GB, more often changes to existing (even old) records

change rate: 2,359,934 file transfers last month (from site A to site B;

change rate on the DB to be found out)