

Searching heterogeneous data-sources: Master thesis problem statement

Vidmantas Zemleris, October 2, 2012

1 Introduction

At large scientific collaborations like the CMS Experiment at CERN's LHC that includes more than 3000 collaborators data usually resides on a fair number of autonomous and heterogeneous proprietary systems each serving it's own purpose ¹. As data stored on one system may be related to data residing on other systems², users are in need of a centralized and easy-to-use solution for locating and combining data from all these multiple services.

Using a highly structured language like SQL is problematic because users need to know not only the language but also where to find the information and also lots of technical details like schema. A data integration system based on simple structured queries is already in place. Still **various improvements including support** for less restricted keyword queries and improvements to system's usability and performance still have to be researched.

2 Case study: the CMS Experiment at CERN

Users' information need may vary greatly depending on their role, however most of the time they are interested in locating full set of entities matching some selection criteria, e.g.:

- find all *files* from *dataset(s)* matching wild-card query each containing some of the 'interesting' *runs* from a list provided (Release validation teams)
- find (all) *datasets* related some specific physics phenomena³ together with conditions describing how this data was recorded by detector or simulated which are present in separate autonomous system than the datasets (Physicists)
- find all *datasets* matching some pattern stored at a given *site* (filtering on entities stored on separate services)

For more use-cases of data retrieval at CMS Experiment see [DGK⁺08].

The Data Aggregation System

The Data Aggregation System (DAS)[KEM10, BKEM11] was created which allows integrated access to a number of proprietary data-sources by processing user's queries on demand - it queries the data-sources, merges the results, and caches them for subsequent uses.

Currently the queries specify what entity the user is interested in (dataset, file, etc) and provide selection criteria (attribute=value, name BETWEEN [v1, v2]) operators. The combined query results could be later 'pipped' for further filtering and aggregation (min, avg, etc), e.g.:

```
dataset=*RelVal* | grep dataset.nevents >1000 | avg(dataset.size), median(dataset.size)
```

The query above would return average and median datasets sizes of ones containing *RelVal* in their name having more than 1000 events.

Queries could be run either from web browser or through command line interface where the results could be fed into another application (e.g. program doing physics analysis or automatic software release validation).

¹For instance, at CERN, due to many reasons (e.g. research and need of freedom, politics of institutes involved) software projects usually evolve in independent fashion resulting in fair number of proprietary systems[KPGM00]. Further high turnover makes it harder extending these systems

²For example, datasets containing physics events are registered at DBS, while the physical location of files is tracked by Phedex which also takes care of their transfers within the worldwide grid storage

³In case of dataset this data is present in filename or run

3 Problem statement

3.1 Ease of use: Query Language over the mediated schema

Even a very simple structured query language that also contain entity names over the mediated schema may seem **hard** to learn, especially in the beginning. On the other hand, at CERN the names in the mediated schema are referring to real-world entities that **are fairly consistently named** (even though there may exist slight differences in their naming on different data-sources). So some sort of guiding helping user to build the query shall be useful.

Further a minimum set of search predicates is imposed by APIs (mainly because of performance reasons) and user has to be at least informed what is he expected to provide.

Also to consider:

- input ambiguity, typos
- **what about ranked search?**

3.2 Performance

3.2.1 Handling distributed search efficiently

3.2.2 Other performance issues

- more intelligent caching needed:
 - given an entity received from provider, determining if it is useful to cache for long-term
 - **At what level are we caching now? query or individual query results?**
 - could make use of even older information warning user
 - could also have different validity dates for certain fields. if certain field is not explicitly requested, even a very old cache could be used.
- query rewriting then used with grep=smf then the same item is available as selection key. check how many instances.
- pre-fetching common (sub-)queries: determining manually and/or automatically
- scale testing - if we are storing lots of historical info. MongoDB is not so performant if DB cant fit in memory. One of well known solutions is installing SSD.
- ? with current API's for aggregation service has to fetch ALL records not just the first page..
- ??? do we have cases then multiple APIs of the same service could be called for the same data, can we choose just one???
- ? showing partial results

4 Preliminary Literature review

TODO: Overview

4.1 Keyword search over DBs

The problem of Keyword search over Relational Databases (or also semi-structured sources like XML) has received a significant attention by the research community over the last decade. There exist many papers and a couple of quite comprehensive PhD. dissertations[Liu11, Cof12] describing performance optimization details, etc.

There was a number of variations of the problem including returning of only Top-k results expecting that users only care about the quality of limited number of results but not completeness of results, ... **describe this in details**

Ranking Query Templates based on keyword query

A simple way to access relational database could be through a set of predefined named query templates (SQL with selection parameters or operators still to be specified) exposed to a user as a Form that the user has to fill in.

[CBC⁺09] proposes alternative approach for processing keyword queries over relational databases: given a keyword query, instead of returning database tuples one could rank query forms that best matches the query for user to choose the right one (if they are properly named this is fairly easy). The ranking is done by **TODO**.

An interesting feature of this approach is that a Query Template is functionally similar to any autonomous web service (which given the parameters would in turn execute that query on its database). In case of the Data Aggregation System, a user after entering a keyword query could be provided with a ranked list of structured queries (attribute=value) that could be processed given data source constraints (e.g. parameters required) and if needed he could refine his search (e.g. provide more parameters).

Semantic approach

Researchers from *ETHZ* have been investigating Generation of SQL for Business users in collaboration with *Credit Suisse*. To converting natural language queries (that in addition to keyword search could convey some semantic structure) into SQL statements they employed some natural language processing and used meta-data describing the schema and the domain (ontologies at multiple levels).

Even on a large data-warehouse of 220GB data with a complex schema of 300 tables they reported fairly promising results that proves that if good meta-data is available, generating even quite complex SQL (n-way joins with aggregations, etc) is quite feasible for computer making it 'much easier for business users to interactively explore highly-complex data warehouses' [BJK⁺12, p.932].

Automatic mapping between distributed services

an interesting approach but not for CERN...

Question Answering

4.2 Searching Heterogeneous web services

There doesn't seem to be much of recent papers going towards this direction (except from specific search systems like flight fare comparisons), however these are based on structured arguments.

In 1990s there were was much research on search based service integration systems, for instance *Information Manifold* following Local-as-View approach and *TSIMMIS* developed at Stanford following Global-as-View approach.

The *Information Manifold* is quite resembling to DAS , so it could be useful checking it's papers. **TODO: finish**

Performance: Answering Queries using Views

Some other approaches could include:

- creating virtual documents offline by joining tables, for instantaneous search results by employing (Indexing Relational Database Content Offline for Efficient Keyword-Based Search, 2003)
- (Efficient Keyword Search Across Heterogeneous Relational Databases, 2007)
 - : combines schema matching and structure discovery techniques to find approximate foreign-key joins

5 Possible approaches

5.1 Ease of use

- javascript interpreter
 - : field + description -- dropdown searchable list
- if entered keyword query suggest a structured query, optionally with description in natural language
- select entity (e.g. dataset) --> give all available keys from forms, however there are still constraints

5.2 Performance

5.2.1 Continuous view maintenance at Data providers with large DBs

Use either *materialized refresh fast views with query rewriting* (Oracle; completely transparent for proprietary apps)[ea11] or some other continuous view maintenance tool (e.g. DBToaster⁴) to improve performance of heavy queries containing joins and/or aggregations.

5.3 Integrating distributed information efficiently

Bloom-join (which could be quite transparent and implemented even on DB side [pure SQL is possible for MySQL, to check for Oracle]- take a query and bit-vector as parameter) , lazy pagination (and order required for aggregation) - this is not yet supported by any of the data-service APIs
integration at DAS level. (at source DB could be more performant a little)

5.4 More intelligent caching

Some of the data entities instances changing very rarely, for example in DBS system old datasets would never change, while new ones are constantly added (still some of their attributes may change, in this example validity).

Cached copy may be shown before hand while up to date results could be retrieved on user's request.

Decide if to cache or not...

An automatic change rate prediction could be useful to efficiently balance between caching and retrieving results.

Also: Caching of intermediary results, not only the queries

6 Work status

TODO

- obtained DB copy of biggest data provider DBS (currently 80 GB + 200 GB indexes)
- preliminary literature review (to be continued deeper + waiting for book arrival: Principles of Data Science)
- some analysis of DAS logs

Upcomming Work items:

- couple of fairly simple prototypes of UI/access patterns for simpler DAS querying
- check on performance improvements after creation of materialized view(s)
- interview (more) DAS users

⁴<http://www.dbtoaster.org> that is being developed at EPFL

References

- [BJK⁺12] Lukas Blunschi, Claudio Jossen, Donald Kossmann, Magdalini Mori, and Kurt Stockinger. Soda: generating sql for business users. *Proc. VLDB Endow.*, 5(10):932–943, June 2012.
- [BKEM11] G Ball, V Kuznetsov, D Evans, and S Metson. Data aggregation system - a system for information retrieval on demand over relational and non-relational distributed data sources. *Journal of Physics: Conference Series*, 331(4):042029, 2011.
- [CBC⁺09] Eric Chu, Akanksha Baid, Xiaoyong Chai, AnHai Doan, and Jeffrey Naughton. Combining keyword search and forms for ad hoc querying of databases. In *Proceedings of the 2009 ACM SIGMOD International Conference on Management of data*, SIGMOD '09, pages 349–360, New York, NY, USA, 2009. ACM.
- [Cof12] Joel Coffman. *Toward Practical Relational Keyword Search Systems*. PhD thesis, May 2012.
- [DGK⁺08] A Dolgert, L Gibbons, V Kuznetsov, C D Jones, and D Riley. A multi-dimensional view on information retrieval of cms data. *Journal of Physics: Conference Series*, 119(7):072013, 2008.
- [ea11] P Lane et al. Oracle database data warehousing guide, 11g release 2 (11.2). chapter 9: Basic materialized views, September 2011.
- [KEM10] Valentin Kuznetsov, Dave Evans, and Simon Metson. The cms data aggregation system. *Procedia Computer Science*, 1(1):1535 – 1543, 2010. [jce:title;ICCS 2010;ce:title](#).
- [KPGM00] Christoph Koch, Paolo Petta, Jean-Marie Le Goff, and Richard McCatchey. On information integration in large scientific collaborations, 2000.
- [Liu11] Ziyang Liu. *Enhancing the usability of complex structured data by supporting keyword searches*. PhD thesis, Tempe, AZ, USA, 2011. AAI3465038.